

STM32 same while loop code but compiled to different assembly code

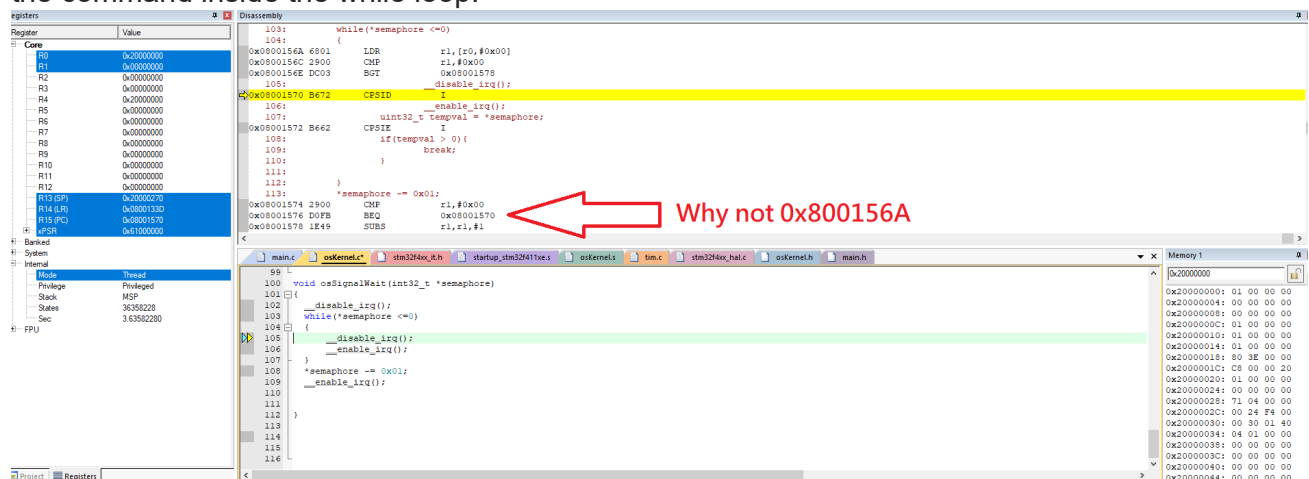
Asked yesterday Active today Viewed 94 times

2

I'm learning RTOS on stm32F411RE board (Cortex-M4). I use MDK uVision v5. I encounter a problem of C code **while** loop. The code in the following is exactly the same in my project and the instructor's project (on Udemy), however, after compiling both project (on my PC), the assembly code look's different. I want to ask what makes this different. Thank you.

```
void osSignalWait(int32_t *semaphore)
{
    __disable_irq();
    while(*semaphore <=0)
    {
        __disable_irq();
        __enable_irq();
    }
    *semaphore -= 0x01;
    __enable_irq();
}
```

In the debug view (see image), if the condition does not match, it does not go to load the real value **LDR r1,[r0,#0x00]** and then do the comparison. Instead, it compares and goes to execute the command inside the while loop.



My code compiled below

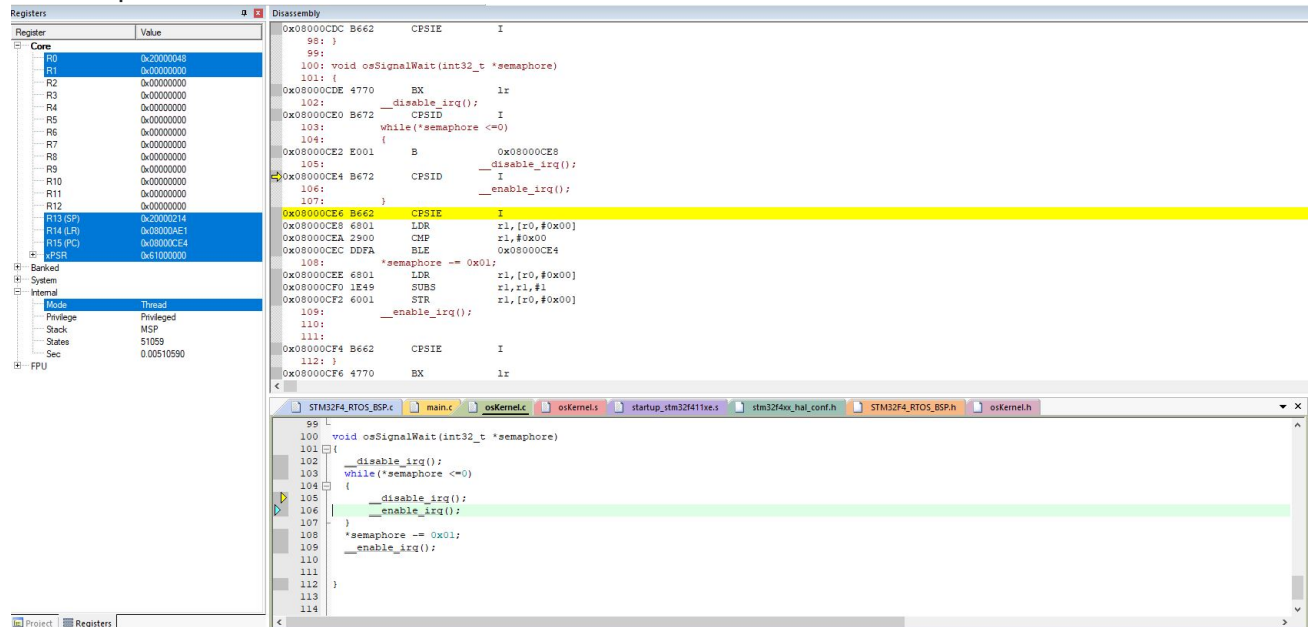
```
100: void osSignalWait(int32_t *semaphore)
101: {
0x08001566 4770    BX            lr
102:    __disable_irq();
103:    while(*semaphore <=0)
104:    {
0x08001568 B672    CPSID        I
101:    {
102:        __disable_irq();
103:        while(*semaphore <=0)
104:        {
0x0800156A 6801    LDR            r1,[r0,#0x00]
0x0800156C E001    B             0x08001572
105:        __disable_irq();
```

```

0x0800156E B672 CPSID I
106:                __enable_irq();
107:            }
108:            *semaphore -= 0x01;
0x08001570 B662 CPSIE I
0x08001572 2900 CMP r1,#0x00
0x08001574 DDFB BLE 0x0800156E
0x08001576 1E49 SUBS r1,r1,#1
109:            __enable_irq();
0x08001578 6001 STR r1,[r0,#0x00]
0x0800157A B662 CPSIE I
110: }

```

If I compile the instructor's (on Udemy) code (on my PC using his project), the assembly code look's different (with exactly the same while loop code). It would load the real value again and do the comparison.



Instructor's code compiled below (Compiled on my PC)

```

100: void osSignalWait(int32_t *semaphore)
101: {
0x08000CDE 4770 BX lr
102:     __disable_irq();
0x08000CE0 B672 CPSID I
103:     while(*semaphore <=0)
104:     {
0x08000CE2 E001 B 0x08000CE8
105:         __disable_irq();
0x08000CE4 B672 CPSID I
106:         __enable_irq();
107:     }
0x08000CE6 B662 CPSIE I
0x08000CE8 6801 LDR r1,[r0,#0x00]
0x08000CEA 2900 CMP r1,#0x00
0x08000CEC DDFB BLE 0x08000CE4
108:     *semaphore -= 0x01;
0x08000CEE 6801 LDR r1,[r0,#0x00]
0x08000CF0 1E49 SUBS r1,r1,#1
0x08000CF2 6001 STR r1,[r0,#0x00]
109:     __enable_irq();
110:
111:

```

```
0x08000CF4 B662      CPSIE      I
112: }
```

c assembly while-loop stm32 cortex-m

edited 22 hours ago

asked yesterday



Dung-Yi

23 ● 4



New contributor

3 ▲ Different compiler version? Different compiler options? – Paul Ogilvie yesterday



▲ Did you ask your instructor? – Jabberwocky yesterday



Hi, because the instructor never reply to students, I have to ask the question here. – Dung-Yi yesterday

@PaulOgilvie Which compiler option do I need to look into? Thank you – Dung-Yi yesterday



OK, so the generated assembly code is not the same. Is this a problem? Does *your* code work as expected? If it's not and if the code works, then go on and assume your instructor uses a different version of the compiler and/or uses different compiler settings. Which settings depends on your compiler, mainly optimisation settings, but who knows... If your instructor doesn't answer, you're out of luck. – Jabberwocky yesterday

@Jabberwocky Thank you. My code doesn't work as expected (If I change the Program Counter Register during the debug mode to the LDR r1,[r0, #0x00] address, the code would work!! I want to know how to make the while loop check the real value (in my assembly code, it doesn't check it) – Dung-Yi yesterday

1 ▲ @Dung-Yi in the instructor's code image you didn't show the very first line of the function. Yes this matters for us as we can't assume anything. – Jabberwocky yesterday



1 ▲ FWIW: I think the body of the while loop should be `__enable_irq(); __disable_irq();` in *that* order. – Chris Hall yesterday ✎



@Jabberwocky Thank you so much. I've updated the image. Please advise. Thank you – Dung-Yi 23 hours ago

2 Answers



2

Since you aren't telling the compiler `semaphore` can change during the execution of this function, your compiler has decided to optimize your code and load the value of semaphore only once and use its copy in the while loop, then only write the result in the end. As it is written now, there's no reason for the compiler to assume this could be harmful.



To notify the compiler a variable *can* change outside the function, during the execution of that function, please use the `volatile` keyword, see:
<https://en.cppreference.com/w/c/language/volatile>



In that case, your code would become:

```
void osSignalWait(volatile int32_t *semaphore)
{
    __disable_irq();
    while(*semaphore <=0)
    {
        __disable_irq();           // Note: I think the order is wrong...
        __enable_irq();
    }
    *semaphore -= 0x01;
    __enable_irq();
}
```

By the way, calling `__disable_irq` twice (once before the while loop, then at the start inside the loop) then `__enable_irq` seems a bit wonky, don't you mean enable (and do something) then disable within the while loop?

edited 6 hours ago

answered yesterday



Elijan9

875 ● 2 ● 13 ● 16

Thank you so much. It works and solves my problem in my project. But the instructor's doesn't use volatile, how can the debugger knows that. (The instructor's code works perfectly) – [Dung-Yi](#) yesterday ✎

▲ Without optimizations on, the compiler may or may not do the load/store for each iteration. A different compiler or different compiler options may produce a different result. – [Elijan9](#) yesterday

2 ▲ @Dung-Yi: the *debugger* doesn't know that. It only reloads because it must have been compiled without optimization, so the compiler treated *everything* kind of like `volatile`. See [Why does clang produce inefficient asm with -O0 \(for this simple floating point sum\)?](#). Or specifically for that bug in your and your instructor's code: [MCU programming - C++ O2 optimization breaks while loop / Multithreading program stuck in optimized mode but runs normally in -O0](#) – [Peter Cordes](#) 23 hours ago

1 ▲ And BTW, the portable C choice for communication between cores / threads is `_Atomic int32_t *` (with `memory_order_relaxed`). (If you're disabling IRQs, you can use a separate atomic load and atomic store, not `-=`). @Dung-Yi: Hmm, I think your code is broken: **the `while()` loop body leaves interrupts enabled before doing `-=`**. I think you want `enable(); disable();` as the loop body to give interrupt handlers a chance to run while you're spin-waiting. Or better, spin with interrupts enabled until you see `*sem <= 0`, then try disable, check, and decrement. – [Peter Cordes](#) 23 hours ago ✎

▲ @Elijan9: I think enable/disable is just there so interrupt handlers can run while spinning, not to actually do anything between them. – [Peter Cordes](#) 23 hours ago

1 ▲ @Elijan9 - The loop body currently disables interrupts and then enables them, which as Peter Cordes mentions, is backwards. The purpose of those in the loop body (when done in the correct order) is to ensure that when the semaphore becomes available, the loop body exits with interrupts disabled. In the meantime, repeatedly enabling them ensures that interrupts can be serviced during the spin-wait. – [phonetagger](#) 23 hours ago

Thank you guys so much. I appreciate your kindness based on my simple and a little bit silly question. (I've done so much research for almost 1 week but cannot find the correct keyword for the answer. I hope the question is not duplicate) – [Dung-Yi](#) 23 hours ago

1 ▲ @Dung-Yi - The spin-wait technique is a very crude way of handling semaphores. In a good preemptive RTOS, you would be able to wait on the semaphore without spinning. Spinning wastes CPU resources (instruction cycles). The OS (or RTOS) should have a way of waiting/pending on the semaphore by calling some OS API function, and when it returns, either you have the semaphore or the wait timed out (if you gave it a timeout). – [phonetagger](#) 23 hours ago

@phonetagger Do you mean using the Cooperative spin-lock? Insert a yield function in the middle of disable(); enable(); I use the following code to do the yield #define ICSR (*(volatile uint32_t *)0xE000ED04) void osThreadYield(void) { SysTick->VAL = 0; ICSR = 0x04000000; // trigger SysTick } – [Dung-Yi](#) 23 hours ago

1 ▲ No volatile is needed here at all. The builtin asm has the `memory` clobber acting as a compiler memory barrier. – [P__J__](#) 23 hours ago

▲ @P__J__ - Though the assembly that results from compiling the OP's posted source may read the value that `semaphore` points to on each loop, it would be legal, if `semaphore` is **not** defined as `volatile`, for the compiler to perform an optimization by emitting instructions that read it only once, before the loop, and keep re-using the value without reading it again on each loop. By defining it as `volatile int32_t *semaphore` instead of `int32_t *semaphore`, you prevent the compiler from making that optimization. A memory barrier by itself doesn't solve that problem. – [phonetagger](#) 23 hours ago

1 ▲ @Dung-Yi I do not mean using a "cooperative spin-lock". I don't know what sort of RTOS you are using; home-brew perhaps? (Did your professor write it?) Any good commercial RTOS would have a "wait on semaphore" API function that puts your thread to sleep until the semaphore is made available by another thread or interrupt. While your thread is sleeping, it consumes no CPU instruction cycles at all. Using a yield or delay function is better than plain spinning, but it still takes intermittent CPU cycles; it's not as good as an API function that's specifically made for waiting on a semaphore. – [phonetagger](#) 22 hours ago

▲ @phonetagger - it is not memory barrier instruction only compiler memory clobber. It is something different – [P__J__](#) 22 hours ago

▲ @P__J__ - Where is this `memory` clobber thing you speak of? In which asm listing? Part of which instruction? – [phonetagger](#) 22 hours ago

▲ the CMSIS functions: `__attribute__((always_inline)) __STATIC_INLINE void __enable_irq(void) { __ASM volatile ("cpsie i" : : : "memory"); }` and disable as well. – [P__J__](#) 22 hours ago

▲ @phonetagger: P__J__ is talking about a GNU C inline asm statement with an empty template string but a clobber list that declares to the compiler that arbitrary vars in memory might be modified. `asm("::: "memory")`. The optimizer treats it like a black-box function call it can't inline, so only non-escaped local vars can stay in regs. This is a barrier against compile-time reordering and load hoisting, but compiles to zero instructions. preshing.com/20120625/memory-ordering-at-compile-time. I don't see any advantage to doing that over using `volatile` or `_Atomic int`, though – [Peter Cordes](#) 22 hours ago

▲ @P__J__ - OK, perhaps it **should** work as you say, but clearly the asm in the OP's first listing shows it doesn't work as you say. According to stackoverflow.com/a/47103378/1245420, either `volatile` or `"memory"` should work, but in this case clearly `"memory"` doesn't. Using `volatile` is the idiomatic (portable) C and C++ way of telling the compiler to avoid optimizing accesses to/from a variable, and I suspect it will work properly for the OP. – [phonetagger](#) 22 hours ago

▲ @phonetagger: It works with GCC 8.2: godbolt.org/z/fbpe2d. Perhaps Keil is buggy? If memory access gets reordered wrt. `enable / __disable_irq()` that's a potential problem, if any code ever depends on that for non-volatile accesses. Also note that GCC avoids a 2nd load of the semaphore with the barriers but not volatile case. It still has the value in a register from the loop condition when compiling `*semaphore -- 1;` so it only needs a `sub` and `str`. To get that with `volatile` or `_Atomic`, you'd need to load into a tmp var as part of the loop condition. – [Peter Cordes](#) 22 hours ago

1 ▲ @PeterCordes: To allow the interrupt handlers to run while spinning, in general you want to use enable then disable in the loop, not the other way around. The original poster uses them in the opposite order and I choose not to change that as well, but I already mentioned in my post this seems a bit wonky and that I think it should be the other way around. It makes more sense to disable the interrupt while checking the loop condition and enable it for a short period within the loop... – [Elijan9](#) 7 hours ago

▲ I wasn't sure you'd understood that enabling interrupts might just be to let an IRQ handler run. I thought you meant "and do something" as in run some code in the loop body that runs with interrupts enabled

while spin-waiting. But now I think you meant "let an IRQ handler do something". And yes, agreed that the only sane order here is `disable; while(){enable;disable;}` – [Peter Cordes](#) 4 hours ago

@Elijan9 Thank you for bringing up this concern. – [Dung-Yi](#) 6 mins ago

@Elijan9 Thank you for bringing up this concern. I put something like yield function by first setting `SysTick->VAL = 0;` and then `((volatile uint32_t *)0xE000ED04) = 0x04000000;` in the middle of disable and enable to trigger the SysTick and perform the context switch – [Dung-Yi](#) just now [Edit](#)

This very well known keil over optimisation bug. Reported many times. Having memory clobber it should read the memory every time.

0

Here is an example how clobbers work

```
#include <stdint.h>

unsigned x;
volatile unsigned y;

int foo()
{
    while(x < 1000);
}

int bar()
{
    while(x < 1000) asm("");
}
```

```
foo:
    ldr    r3, .L5
    ldr    r3, [r3]
    cmp    r3, #1000
    bxcs   lr
.L3:
    b      .L3
.L5:
    .word  x
bar:
    ldr    r1, .L11
    ldr    r2, .L11+4
    ldr    r3, [r1]
    cmp    r3, r2
    bxhi   lr
.L9:
    ldr    r3, [r1]
    cmp    r3, r2
    bls    .L9
    bx     lr
.L11:
    .word  x
    .word  999
```

answered 22 hours ago



[P_J](#)

20.5k ● 3 ● 12 ● 36

- 1 ▲ That's not an over-optimization bug, it's data-race UB in the C source that you can work around with a memory barrier or volatile. Or actually *avoid* UB with `_Atomic` (with `mo_relaxed` so it can compile to the same asm). Why do you say a barrier is better than `volatile`, or `_Atomic int` (with `mo_relaxed`)? Since this is for a semaphore object that's *only* ever used for synchronization so you never want to let the compiler keep it in a register across multiple reads. `volatile` ensures we don't have [invented loads](#) from one read. – Peter Cordes 22 hours ago

Thank you for point out the issue. I did find this bug reported in stack overflow before but never realize that my problem is caused by this!! Thank you. – Dung-Yi 22 hours ago

- ▲ @PeterCordes I do not say what is better. They are not identical and have different uses
 godbolt.org/z/aXSMB – P__J__ 22 hours ago

- ▲ @Dung-Yi: Just to be clear; it's not a *compiler* bug, it's a bug in your code. The compiler is working as intended and optimizing variables into registers under the assumption that no other thread can modify them, unless you tell it otherwise. It's allowed to assume that for non-`_Atomic` and non-`volatile` variables because [otherwise that would be data-race UB](#). If compilers didn't do this, any code that used global variables or even pointers would run slowly, storing after every change and reloading before every use. – Peter Cordes 22 hours ago

- ▲ @P__J__: I think I was getting mixed up between the comment thread on the other answer vs. this. But why post it when there's already an answer suggesting `volatile`? The interesting point is that the "memory" barrier that should be part of the definition of `__disable_irq()` should have already blocked reordering, like it does on GCC. [godbolt.org/z/fbpe2d](#). (But yes I'm aware that memory barriers don't affect non-escaped local vars, like a function arg.) – Peter Cordes 22 hours ago

@PeterCordes Thank you for correction. – Dung-Yi 21 hours ago

- ▲ @PeterCordes just to show particular Keil problem as it is a very specific question – P__J__ 21 hours ago

- ▲ But what problem does your answer show? Both loops use `x` which is *not* `volatile`. GCC and clang would also make asm something like that for that source: `if(x) infloop()` for the no-barrier no-volatile version (because nothing hides the C data-race UB from the compiler), or volatile-equivalent for the version with a barrier. – Peter Cordes 21 hours ago

- ▲ And if this answer works for keil, why doesn't the "memory" clobber in `__enable_irq()` have the same effect? Or is it maybe defined incorrectly in some versions of that header? – Peter Cordes 21 hours ago