

STM32 same while loop code but compiled to different assembly code

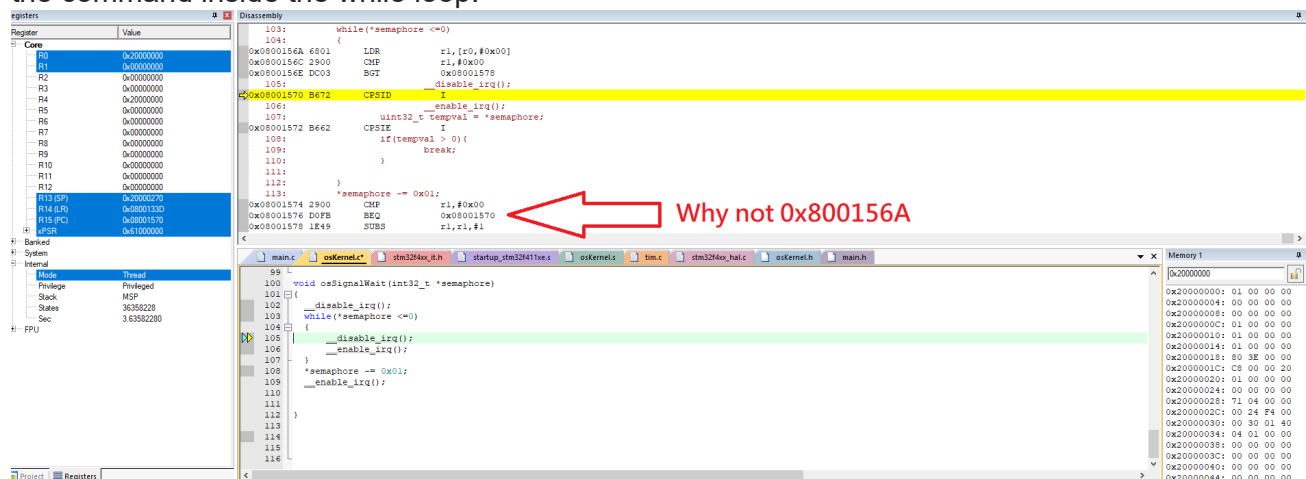
Asked today Active today Viewed 66 times

2

I'm learning RTOS on stm32F411RE board (Cortex-M4). I use MDK uVision v5. I encounter a problem of C code **while** loop. The code in the following is exactly the same in my project and the instructor's project (on Udemy), however, after compiling both project (on my PC), the assembly code look's different. I want to ask what makes this different. Thank you.

```
void osSignalWait(int32_t *semaphore)
{
    __disable_irq();
    while(*semaphore <=0)
    {
        __disable_irq();
        __enable_irq();
    }
    *semaphore -= 0x01;
    __enable_irq();
}
```

In the debug view (see image), if the condition does not match, it does not go to load the real value **LDR r1,[r0,#0x00]** and then do the comparison. Instead, it compares and goes to execute the command inside the while loop.



My code compiled below

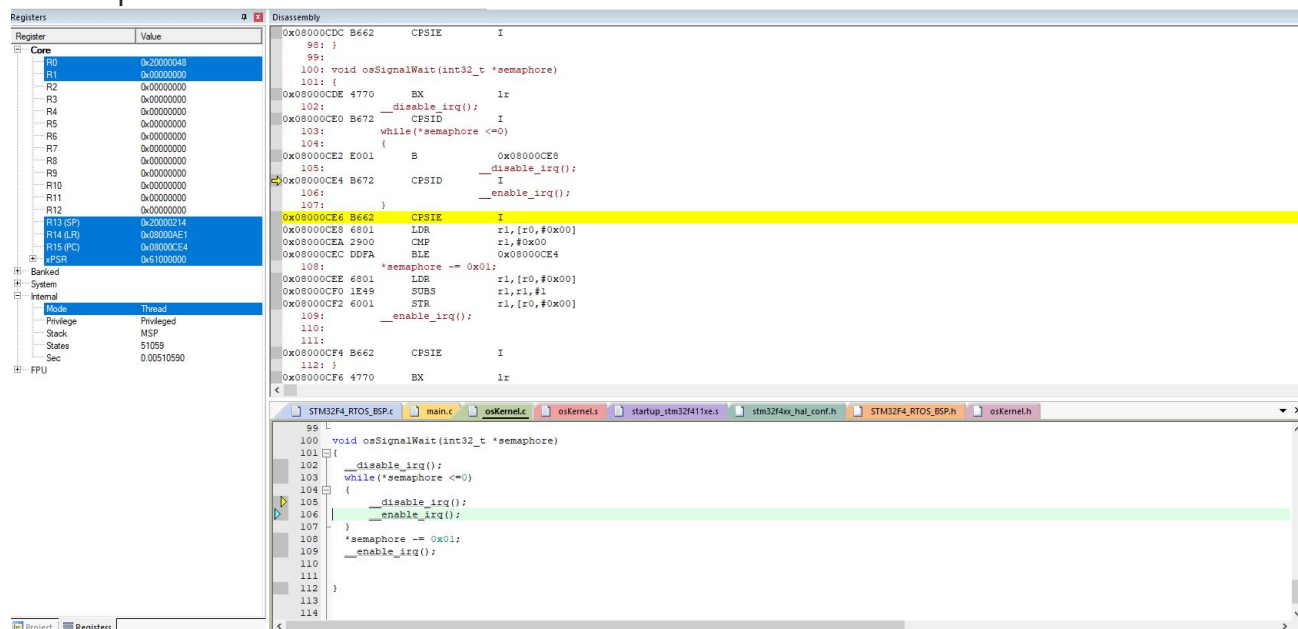
```
100: void osSignalWait(int32_t *semaphore)
101: {
0x08001566 4770    BX            lr
102:    __disable_irq();
103:    while(*semaphore <=0)
104:    {
0x08001568 B672    CPSID        I
101:    {
102:        __disable_irq();
103:        while(*semaphore <=0)
104:        {
0x0800156A 6801    LDR            r1,[r0,#0x00]
0x0800156C E001    B             0x08001572
105:        __disable_irq();
```

```

0x0800156E B672 CPSID I
106:                __enable_irq();
107:            }
108:            *semaphore -= 0x01;
0x08001570 B662 CPSIE I
0x08001572 2900 CMP r1,#0x00
0x08001574 DDFB BLE 0x0800156E
0x08001576 1E49 SUBS r1,r1,#1
109:            __enable_irq();
0x08001578 6001 STR r1,[r0,#0x00]
0x0800157A B662 CPSIE I
110: }

```

If I compile the instructor's (on Udemy) code (on my PC using his project), the assembly code look's different (with exactly the same while loop code). It would load the real value again and do the comparison.



Instructor's code compiled below (Compiled on my PC)

```

100: void osSignalWait(int32_t *semaphore)
101: {
0x08000CDE 4770 BX lr
102:     __disable_irq();
0x08000CE0 B672 CPSID I
103:     while(*semaphore <=0)
104:     {
0x08000CE2 E001 B 0x08000CE8
105:         __disable_irq();
0x08000CE4 B672 CPSID I
106:         __enable_irq();
107:     }
0x08000CE6 B662 CPSIE I
0x08000CE8 6801 LDR r1,[r0,#0x00]
0x08000CEA 2900 CMP r1,#0x00
0x08000CEC DDFB BLE 0x08000CE4
108:     *semaphore -= 0x01;
0x08000CEE 6801 LDR r1,[r0,#0x00]
0x08000CF0 1E49 SUBS r1,r1,#1
0x08000CF2 6001 STR r1,[r0,#0x00]
109:     __enable_irq();
110:
111:

```

```
0x08000CF4 B662 CPSIE I
112: }
```

c assembly while-loop stm32 cortex-m

edited 16 mins ago

asked 2 hours ago



Dung-Yi

23 ● 4



New contributor

2 ▲ Different compiler version? Different compiler options? – Paul Ogilvie 2 hours ago



▲ Did you ask your instructor? – Jabberwocky 2 hours ago



Hi, because the instructor never reply to students, I have to ask the question here. – Dung-Yi 2 hours ago

1 ▲ @Dung-Yi in the instructor's code image you didn't show the very first line of the function. Yes this matters for us as we can't assume anything. – Jabberwocky 2 hours ago



1 ▲ FWIW: I think the body of the while loop should be `__enable_irq(); __disable_irq();` in *that* order. – Chris Hall 2 hours ago ✎



|

2 Answers



2



Since you aren't telling the compiler `semaphore` can change during the execution of this function, your compiler has decided to optimize your code and load the value of semaphore only once and use its copy in the while loop, then only write the result in the end. As it is written now, there's no reason for the compiler to assume this could be harmful.

To notify the compiler a variable *can* change outside the function, during the execution of that function, please use the `volatile` keyword, see:

<https://en.cppreference.com/w/c/language/volatile>

In that case, your code would become:

```
void osSignalWait(volatile int32_t *semaphore)
{
    __disable_irq();
    while(*semaphore <= 0)
    {
        __disable_irq();
        __enable_irq();
    }
    *semaphore -= 0x01;
    __enable_irq();
}
```

By the way, calling `__disable_irq` twice (once before the while loop, then at the start inside the loop) then `__enable_irq` seems a bit wonky, don't you mean enable (and do something) then disable within the while loop?

edited 2 hours ago

answered 2 hours ago



Elijan9

875 ● 2 ● 13 ● 16

- 1 ▲ @Dung-Yi: the *debugger* doesn't know that. It only reloads because it must have been compiled without optimization, so the compiler treated *everything* kind of like `volatile`. See [Why does clang produce inefficient asm with -O0 \(for this simple floating point sum\)?](#). Or specifically for that bug in your and your instructor's code: [MCU programming - C++ O2 optimization breaks while loop / Multithreading program stuck in optimized mode but runs normally in -O0](#) – Peter Cordes 1 hour ago
- 1 ▲ And BTW, the portable C choice for communication between cores / threads is `_Atomic int32_t *` (with `memory_order_relaxed`). (If you're disabling IRQs, you can use a separate atomic load and atomic store, not `--`). @Dung-Yi: Hmm, I think your code is broken: **the `while()` loop body leaves interrupts enabled before doing `--`**. I think you want `enable(); disable();` as the loop body to give interrupt handlers a chance to run while you're spin-waiting. Or better, spin with interrupts enabled until you see `*sem <= 0`, then try disable, check, and decrement. – Peter Cordes 1 hour ago
- 1 ▲ @Elijan9 - The loop body currently disables interrupts and then enables them, which as Peter Cordes mentions, is backwards. The purpose of those in the loop body (when done in the correct order) is to ensure that when the semaphore becomes available, the loop body exits with interrupts disabled. In the meantime, repeatedly enabling them ensures that interrupts can be serviced during the spin-wait. – phonetagger 1 hour ago
- 1 ▲ @Dung-Yi - The spin-wait technique is a very crude way of handling semaphores. In a good preemptive RTOS, you would be able to wait on the semaphore without spinning. Spinning wastes CPU resources (instruction cycles). The OS (or RTOS) should have a way of waiting/pending on the semaphore by calling some OS API function, and when it returns, either you have the semaphore or the wait timed out (if you gave it a timeout). – phonetagger 1 hour ago
- 1 ▲ No volatile is needed here at all. The builtin asm has the `memory` clobber acting as a compiler memory barrier. – P__J__ 58 mins ago

|



0



This very well known keil over optimisation bug. Reported many times. Having memory clobber it should read the memory every time.

Here is an example how clobbers work

```
#include <stdint.h>

unsigned x;
volatile unsigned y;

int foo()
{
    while(x < 1000);
}

int bar()
{
    while(x < 1000) asm(""); // "memory";
}
```

```

foo:
    ldr    r3, .L5
    ldr    r3, [r3]
    cmp    r3, #1000
    bxcs   lr
.L3:
    b      .L3
.L5:
    .word  x
bar:
    ldr    r1, .L11
    ldr    r2, .L11+4
    ldr    r3, [r1]
    cmp    r3, r2
    bxhi   lr
.L9:
    ldr    r3, [r1]
    cmp    r3, r2
    bls    .L9
    bx     lr
.L11:
    .word  x
    .word  999

```

answered 38 mins ago



P_J_

20.4k ● 3 ● 12 ● 36

- ▲ That's not an over-optimization bug, it's data-race UB in the C source that you can work around with a memory barrier or volatile. Or actually *avoid* UB with `_Atomic` (with `mo_relaxed` so it can compile to the same asm). Why do you say a barrier is better than `volatile`, or `_Atomic int` (with `mo_relaxed`)? Since this is for a semaphore object that's *only* ever used for synchronization so you never want to let the compiler keep it in a register across multiple reads. `volatile` ensures we don't have [invented loads](#) from one read. – Peter Cordes 26 mins ago

Thank you for point out the issue. I did find this bug reported in stack overflow before but never realize that my problem is caused by this!! Thank you. – Dung-Yi 19 mins ago

- ▲ @PeterCordes I do not say what is better. They are not identical and have different uses
 ▲ godbolt.org/z/aXSMB – P_J_ 8 mins ago

- ▲ @Dung-Yi: Just to be clear; it's not a *compiler* bug, it's a bug in your code. The compiler is working as intended and optimizing variables into registers under the assumption that no other thread can modify them, unless you tell it otherwise. It's allowed to assume that for non-`_Atomic` and non-`volatile` variables because [otherwise that would be data-race UB](#). If compilers didn't do this, any code that used global variables or even pointers would run slowly, storing after every change and reloading before every use. – Peter Cordes 7 mins ago

- ▲ @P_J_: I think I was getting mixed up between the comment thread on the other answer vs. this. But why post it when there's already an answer suggesting `volatile`? The interesting point is that the "memory" barrier that should be part of the definition of `__disable_irq()` should have already blocked reordering, like it does on GCC. godbolt.org/z/fbpe2d. (But yes I'm aware that memory barriers don't affect non-escaped local vars, like a function arg.) – Peter Cordes 2 mins ago